DynLang-SLAM: Tracking Pipeline Tuning Report

Systematic Ablation Study on Replica Dataset

Ankur Guruprasad
EEE 515 – Computer Vision, Arizona State University
agurupr6@asu.edu

April 2026

1 Introduction

This report documents the systematic parameter tuning and algorithmic improvements applied to the tracking pipeline of DynLang-SLAM, a 3D Gaussian Splatting-based visual SLAM system. All experiments were conducted on the Replica dataset [1] using three scenes (`room0`, `room1`, `room2`) with 200 frames each. The primary metric is Absolute Trajectory Error (ATE RMSE) in centimeters.

1.1 System Overview

The tracking pipeline optimizes a 6-DOF camera pose per frame by minimizing a weighted combination of photometric (L1 + SSIM) and geometric (depth L1) losses against a differentiable Gaussian splatting map rendered via `gsplat` [2]. The pose is parameterized as a quaternion (4D) + translation (3D) and optimized with Adam for a fixed number of iterations.

1.2 Baseline Configuration

- Tracking iterations: 50
- Pose learning rate: 0.002 (constant)
- Loss weights: RGB=0.5, Depth=1.0, SSIM=0.2
- Render downscale: 1 (full resolution 680×1200)
- Keyframe interval: every 5 frames
- Densification: gradient threshold 0.0002, interval 10

2 Initial Baseline Results

Table 1 shows the original ATE values before any tuning.

Table 1: Initial baseline ATE (RMSE, cm) on Replica scenes.

Scene	ATE (cm)	PSNR	Target
room0	0.72	–	<1.0
room1	10.34	–	<5.0
room2	18.82	–	<5.0

3 Methods Explored

3.1 Depth Weight Tuning

We varied the depth loss weight to find the optimal balance between geometric and photometric alignment.

Table 2: Depth weight ablation (ATE RMSE in cm).

depth_w	room0	room1	room2
1.0 (baseline)	1.67	9.79	4.43
1.2	1.43	8.15	3.73
1.5	–	9.50	1.77

Finding: `depth_w=1.2` is the best universal compromise. Higher values (1.5) dramatically help room2 but catastrophically hurt room1, suggesting scene-dependent sensitivity to depth weighting.

3.2 Velocity Initialization

Constant velocity motion model: the initial pose for frame t is extrapolated as $\hat{T}_t = T_{t-1} \cdot (T_{t-2}^{-1} \cdot T_{t-1})$.

Table 3: Velocity initialization ablation.

Config	room0	room1
Without velocity init	1.51	15.88
With velocity init	1.43	10.98

Finding: Velocity initialization is critical for room1 (31% improvement) where camera motion is aggressive. Neutral for room0 where motion is smooth.

3.3 Isotropic Regularization Weight

Controls the strength of isotropic shape regularization on Gaussian primitives during mapping.

Finding: Stronger isotropic regularization (`iso=0.05`) helps room0 by producing rounder, more stable Gaussians

Table 4: Isotropic regularization ablation on room0.

iso_reg_weight	room0 ATE (cm)
0.01	1.70
0.05	1.43

that provide a cleaner tracking signal.

3.4 Densification

Tested the effect of removing adaptive Gaussian densification entirely.

Table 5: Densification ablation on room0.

Config	room0 ATE (cm)
With densification	1.51
Without densification	2.33

Finding: Densification is essential – without it, the map becomes too sparse for accurate tracking.

3.5 Motion Magnitude Clamping

Clamped the predicted velocity-based motion to prevent extreme extrapolations during sudden camera jerks (5× running average, 0.05m minimum).

Table 6: Motion clamping ablation.

Config	room0	room1
Without clamping	1.67	9.79
With clamping	1.78	7.54

Finding: Clamping helped room1 (23% better) but slightly hurt room0. The tradeoff was not favorable, so this was **removed**.

3.6 Outlier Rejection with Re-Tracking

Detected tracking failures via loss thresholds and re-initialized from the previous frame's pose.

Finding: Catastrophic failure. Re-tracking from the previous pose disrupted the velocity model's accumulated motion estimates, causing cascading drift. **Removed.**

4 Failed Approaches from Literature

Based on techniques from GS-SLAM, TVG-SLAM, CG-SLAM, and others, we implemented and tested several advanced methods. All were evaluated against the tuned baseline (`depth_w=1.2`, velocity init, `iso=0.05`).

Table 7: Outlier rejection ablation.

Config	room1 ATE (cm)
Without outlier rejection	9.79
With outlier rejection	26.14

4.1 Coarse-to-Fine Tracking

Inspired by GS-SLAM [3]: render at lower resolution for early iterations (coarse pose estimation), then refine at full resolution.

Table 8: Coarse-to-fine tracking ablation on room0.

Config	room0 ATE (cm)
Baseline (no C2F)	1.43
C2F: 4× ds, 60% coarse	4.07
C2F: 2× ds, 40% coarse	3.41*

*Combined with other changes (LR scheduling, Huber loss, depth-aware scales).

Finding: Significant regression. Indoor scenes rely on high-frequency geometric details (chair legs, desk edges). Downscaling destroys the alignment signal that makes the optimizer converge. Unlike outdoor/large-scale scenes where coarse-to-fine helps, Replica rooms require full-resolution tracking throughout.

4.2 Learning Rate Scheduling

Applied CosineAnnealingLR to the pose optimizer, decaying from 0.002 to 0.0006 over 50 iterations.

Table 9: LR scheduling ablation on room0.

Config	room0 ATE (cm)
Constant LR (0.002)	1.77
CosineAnnealingLR + Huber	7.03

Finding: Severe regression. With only 50 iterations, decaying the learning rate traps the pose in local minima before it can traverse the loss landscape. Constant LR is essential for this short optimization horizon.

4.3 Huber Depth Loss

Replaced L1 depth loss with Huber loss ($\delta = 0.1$m) for robustness to depth outliers.

Finding: Contributed to regression (tested jointly with LR scheduling). Large depth errors are often the *only* valid signal pulling a wrong pose back to the correct basin of attraction. Clipping these gradients (as Huber does for $|e| > \delta$) prevents convergence.

4.4 Depth-Aware Scale Initialization

Initialized Gaussian scales proportional to depth × pixel footprint instead of a uniform 1cm.

Finding: Caused 30% fewer Gaussians (69K vs 99K). Larger initial scales meant fewer pixels were flagged as "unmapped" during expansion, resulting in a sparser map that hurt tracking.

5 Successful Method: Dynamic Loss Weighting

5.1 Motivation

Instead of changing the learning rate (which failed), we change the **loss weight distribution** across tracking iterations. The key insight: depth loss creates a convex basin of attraction for coarse alignment, while RGB/SSIM losses provide high-frequency signal for sub-pixel refinement.

5.2 Method

Given base loss weights w_{rgb}, w_{depth}, w_{ssim}, we interpolate per-iteration weights using a linear schedule parameterized by $t = i/(N-1) \in [0, 1]$:

$$s_{\text{depth}}(t) = 2.0 - 1.5t \qquad (2.0 \to 0.5) \tag{1}$$

$$s_{\text{rgb}}(t) = 0.3 + 1.2t \qquad (0.3 \to 1.5) \tag{2}$$

The effective weights at iteration i are:

$$w_{\text{depth}}^{(i)} = w_{\text{depth}} \cdot s_{\text{depth}}(t) \tag{3}$$

$$w_{\text{rgb}}^{(i)} = w_{\text{rgb}} \cdot s_{\text{rgb}}(t) \tag{4}$$

$$w_{\text{ssim}}^{(i)} = w_{\text{ssim}} \cdot s_{\text{rgb}}(t) \tag{5}$$

This keeps the total gradient magnitude roughly constant (unlike LR scheduling) while steering the gradient *direction* from geometry-dominant to appearance-dominant.

5.3 Results

Table 10: Dynamic loss weighting results (ATE RMSE in cm). All scenes improved.

Scene	Original	Tuned BL	Dyn. Wt.	Δ
room0	0.72	1.43	**1.12**	-22%
room1	10.34	9.79	**3.32**	-66%
room2	18.82	3.73	**1.50**	-60%

5.4 Analysis

Dynamic loss weighting achieves the same goal as coarse-to-fine tracking (broad convergence first, then refinement) but through the loss landscape rather than the image resolution:

- **Early iterations (depth-heavy):** The depth loss provides a smooth, convex error surface that guides the optimizer into the correct basin of attraction, even when the initial pose estimate is far from ground truth.
- **Late iterations (RGB/SSIM-heavy):** Once the pose is approximately correct, photometric losses exploit high-frequency texture and edge information for sub-pixel alignment that depth sensors cannot provide.

The improvement is especially dramatic for room1 (66% reduction), which features aggressive camera motion. The depth-heavy initialization phase prevents the catastrophic drift that previously occurred around frames 100–150.

6 Other Improvements

Two additional low-risk changes were retained:

- **Quaternion normalization:** Normalizing the quaternion after each Adam step prevents the quaternion magnitude from drifting, which would otherwise distort Adam's effective learning rate.
- **Fast SE(3) inverse:** Replacing `torch.inverse()` with the analytical rigid-body inverse $T^{-1} = [R^\top | -R^\top t]$ in non-differentiable code paths. Mathematically identical, computationally cheaper.

7 Summary of All Experiments

Table 11 provides a comprehensive overview of every method tested.

8 Key Lessons Learned

1. **One change at a time.** Combining multiple modifications simultaneously caused unpredictable interactions and made debugging impossible. Sequential ablation is essential.
2. **Pose optimization is short and non-convex.** With only 40–50 iterations, techniques that work for long training runs (LR scheduling, Huber loss) can be counterproductive.
3. **Depth is a convex anchor.** Heavy depth weighting in early iterations creates a smooth loss landscape that prevents local minima entrapment.

Table 11: Complete ablation summary. All values are ATE RMSE (cm) on 200 frames. Green = improvement, Red = regression.

#	Method	room0	room1	room2	Verdict
0	Original baseline	0.72	10.34	18.82	Starting point
1	+ depth_w=1.2	1.43	8.15	3.73	Kept
2	+ Velocity init	1.43	10.98†	–	Kept
3	+ iso_reg=0.05	1.43	–	–	Kept
4	+ Motion clamping	1.78	7.54	–	Removed
5	+ Outlier rejection	–	26.14	–	Removed
6	+ Coarse-to-fine (4×)	4.07	–	–	Removed
7	+ All changes combined	3.41	–	–	Removed
8	+ LR scheduling + Huber	7.03	–	–	Removed
9	+ Quat norm + fast SE(3)	1.77	–	–	Neutral
10	+ **Dynamic loss weighting**	**1.12**	**3.32**	**1.50**	**Best**

†Without velocity init, room1 = 15.88cm.

4. **Resolution matters for indoor scenes.** Coarse-to-fine tracking fails because indoor environments rely on high-frequency geometric details for alignment.

5. **Don't fight the velocity model.** Outlier rejection with re-tracking disrupted the constant velocity assumption, causing cascading drift.

9 Next Steps

- Full 2000-frame evaluation on ASU Sol supercomputer (A100 GPU)
- Push room0 below 1.0cm target (information-driven keyframe selection)
- Language pipeline: CLIP + SAM + autoencoder feature fusion
- Dynamic object masking: YOLOv8 detection and exclusion

References

[1] Straub et al., "The Replica Dataset: A Digital Replica of Indoor Spaces," *arXiv:1906.05797*, 2019.

[2] Ye et al., "gsplat: An Open-Source Library for Gaussian Splatting," *arXiv:2409.06765*, 2024.

[3] Yan et al., "GS-SLAM: Dense Visual SLAM with 3D Gaussian Splatting," *CVPR*, 2024.